

Data Engineering in the Cloud

Apache Spark - Part I

Xuemao Zhang
East Stroudsburg University

January 18, 2025

Outline

- Introduction to Python
- Introduction to Scala
- Apache Spark
- Data Frame

Python

Python is a high-level, general-purpose programming language.

- Widely used in various domains such as web development, automation, scientific computing, data analysis, and machine learning.
- Rich Ecosystem of Libraries:
 - ▶ Data Manipulation: Libraries like Pandas and NumPy for efficient data handling.
 - ▶ Data Visualization: Tools like Matplotlib, Seaborn, and Plotly for creating insightful visualizations.
 - ▶ Machine Learning: Comprehensive libraries such as Scikit-Learn, TensorFlow, and PyTorch for building and deploying machine learning models.

Python

- You **do not** need to **download** and install python on your computer - We'll use cloud computing [Google Colaboratory](https://colab.research.google.com/notebooks/)
<https://colab.research.google.com/notebooks/>
- Python has no command for declaring a variable. A variable is created the moment you first assign a value to it.
- Variable names are case-sensitive.

```
a = 4  
A = 40  
print(a)
```

```
## 4
```

```
print(A)
```

```
## 40
```

Python

- The Python `print()` function is often used to output variables.
- In the `print()` function, you output multiple variables, separated by a comma:

```
x= y= z = 4
```

```
print(x, y, z)
```

```
## 4 4 4
```

```
# f-string  
print(f"The value of variable x is {x}.")
```

```
## The value of variable x is 4.
```

- We use function `type()` to check data type

```
type(x)
```

```
## <class 'int'>
```

Python

- A Python **list** is an *ordered mutable/changeable* array. Lists *allow duplicate elements* regardless of their type.
- Create a list in Python by using square brackets, separating individual elements with a comma.
 - ▶ Each element can be of any data type.
- All lists in Python are *zero-based indexed*. You can access individual list elements.

```
A = [1, 2, 3, 3.4]
print(A[0],A[1],A[2],A[3])
```

```
## 1 2 3 3.4
```

```
print(A[0:2]) #stop=2; the end element is not included
```

```
## [1, 2]
```

Python functions

- A function takes a list of argument values, performs a computation with those values, and returns a single result. Python gives you many built-in functions.
 - ▶ see [Python Built-in Functions](https://docs.python.org/3/library/functions.html)
<https://docs.python.org/3/library/functions.html>
- We can also create our own functions. These functions are called user-defined functions.
- Functions are declared using the **def** keyword, and the value produced is returned using the **return** keyword. Consider a simple function which returns the square of the input, $y = x^2$.
 - ▶ colon : is used to represent an indented block. It is not for slicing.
 - ▶ Python uses **indentation** to indicate a block of code.
 - ▶ The number of spaces is up to you as a programmer, but it has to be at least one.

Python functions

```
def square(x):  
    return x**2
```

```
x = 2  
y = square(x) # Call the function  
print(x,y)
```

```
## 2 4
```


Scala

- **Scala** is a strong statically typed high-level general-purpose programming language.
 - ▶ Statically typed, meaning types are checked at compile-time.
 - ▶ It is designed to be a scalable language, combining functional and object-oriented programming in a concise, high-performance language.
- Scala is object-oriented, and uses a syntax termed curly-brace which is similar to the language C.
- Typically faster than Python because it compiles to JVM bytecode and runs on the Java Virtual Machine (JVM).
 - ▶ It is often used for scalable systems, big data processing (e.g., with Apache Spark), and back-end services.

Scala

- To avoid installation, we use replit (<https://replit.com/>) to run Scala code.
- Sign up an account and then **Create a new repl** with selecting Scala as the language.
- In the shell, type `scala`
- At the prompt `scala>`, type `println("Hello, Scala!")`

Scala basic syntax

- Case Sensitivity - Scala is case-sensitive
- Class Names - For all class names, the first letter should be in Upper Case.
- Method Names - All method names should start with a Lower Case letter.
- Program File Name - Name of the program file should exactly match the object name. When saving the file you should save it using the object name and append `.scala` to the end of the name.
- `def main(args: Array[String])` - Scala program processing starts from the `main()` method which is a mandatory part of every Scala Program.

Scala variables

- [Data types](#) see the documentation
- Variable declaration: declared using the keyword `val`
- In the following, `String` is the data type

```
val var1: String = "Foo"
```

```
println(var1)
```

- Variable type inference when an initial value is assigned

```
val var2 = 10  
val var3 = "Hello, Scala!"
```

```
println(var3.getClass)
```

Scala variables

- Use `print()` when you want to output multiple pieces of data on the same line
- Use `println()` when you want to display each piece of output on a new line.

```
print(var1,var2,var3)
```

Scala functions

- A Scala function definition has the form

```
def functionName ([list of parameters with data types]) : [return type]  
  function body  
  return [expr]  
}
```

```
def square(x: Double): Double= {  
  val result= x * x  
  return result  
}
```

- or use a shorter format

```
def square(x: Double): Double = x * x
```

Scala functions

- Usage

```
println(square(2))
```

```
val number: Double = 5.1
```

```
val result: Double = square(number)
```

```
println(s"The square of $number is $result")
```

- s string interpolator above allows embedding variables and expressions inside strings.
- f Interpolator: Allows formatted strings using format specifiers, similar to printf in other languages.

```
println(f"The square of $number is $result%.3f")
```

Apache Spark

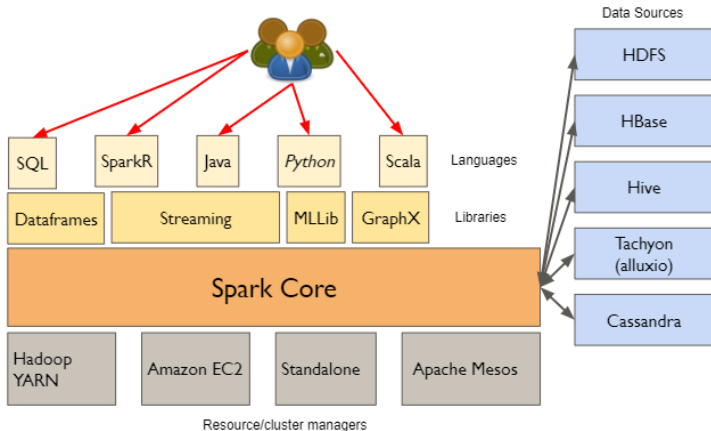
- **Apache Spark** is an open-source unified analytics engine for large-scale data processing.
 - ▶ Spark provides an interface for programming clusters with implicit data parallelism and fault tolerance.
 - ▶ Apache Spark is an **in-memory cluster computing framework** designed to handle a wide range of big data workloads.
 - ▶ Spark was developed to overcome some limitations of *Hadoop's MapReduce*, such as slow processing speed and the complexity of writing MapReduce jobs.
 - ▶ Spark can run on Hadoop YARN and use HDFS for storage.

Apache Spark key features

- In-Memory Computing: Spark performs computations in memory, improving the speed of data processing.
- Cluster Computing Framework: Designed to handle big data workloads across distributed clusters.
- Wide Range of Workloads: Supports various data processing tasks, including:
 - ▶ Batch Processing
 - ▶ Interactive Queries
 - ▶ Streaming Data
 - ▶ Machine Learning
 - ▶ Graph Processing

Components of Apache Spark

Spark Architecture




Core Components of Apache Spark

- Spark Core: The foundation of the Spark framework, providing basic functionality such as task scheduling, memory management, and fault recovery.
- Spark SQL: Allows querying of structured data via SQL, integrating seamlessly with Apache Hive.
- Spark Streaming: Enables real-time data processing and streaming analytics.
- MLlib: A scalable machine learning library offering a variety of algorithms and utilities.
- GraphX: A library for graph processing and analysis.

Spark RDDs


- The processing in distributed computing occurs in numerous phases, with the output of the first stage becoming input to the second stage and the data being shared between the two.



- This kind of processing involves a lot of input output overhead and makes the overall computation slower.
 - ▶ Reduction of the input output operations that slow the computation can be accomplished through in memory data sharing.
 - ▶ The data transferred between in memory is 10-100 times faster than data sharing through a network or a disc.

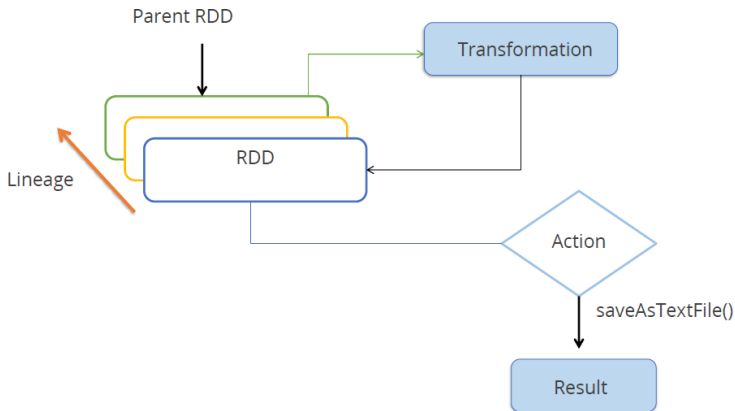
Spark RDDs

- The ideal solution is RDDs which are core concept in Spark.
 - ▶ RDD stands for Resilient Distributed Dataset.
 - ▶ RDD enables fault tolerant, distributed in memory computations.
- An RDD is a collection of objects that are distributed across nodes in a cluster. RDD supports lazy evaluation.



Spark RDDs

- RDD supports lazy evaluation. Lazy evaluation in Spark is a mechanism where the execution of an action will not begin until the action is initiated. Lazy evaluation occurs during Spark transformation.
 - ▶ Lazy evaluation occurs during Spark transformation.




Spark RDDs

- Different Ways to Create Spark RDD

Methods	Ways to use the method	Example
Text File	File or set of files Local or Distributed	<code>textFile = sc.textFile("README.md")</code>
Parallelized Collection	Memory	<code>input = sc.parallelize((1, 2, 3, 4))</code>
From existing RDD	Another RDD	<code>newRdd = input.map(lambda e : (e, 1))</code>
From external datasets	Referencing a dataset in the external storage system	<code>data = spark.read \ .format("org.apache.spark.sql.cassandra")\ .options(table="test", keyspace="test") \ .load()</code>


RDD operations

- Once an RDD is created, two types of operations can be performed
 - ▶ Actions: Act on the entire RDD
 - ▶ Transformations: Create a new RDD from an existing RDD by applying a certain logic
- Actions is an RDD operation that instructs Spark to perform computations and return the results back to the driver. It will be executed only when an action is encountered.




RDD operations

- There are two types of transformation:
 - ▶ Narrow Transformation
 - ▶ Wide Transformation
- Narrow transformations are for data from a single partition only.
 - ▶ Narrow transformations include `map()`, `mapPartition()`, `flatMap()`, `filter()`, and `union()` functions.



RDD operations

- Wide transformation are for data that can be from multiple partitions.
 - ▶ It is the result of GroupByKey() and ReduceByKey() functions
 - ▶ Wide transformations include groupByKey(), aggregateByKey(), aggregate(), join(), and repartition()



Resources for Learning Apache Spark


- [Apache Spark Documentation](https://spark.apache.org/documentation.html) <https://spark.apache.org/documentation.html>
- Book [Learning Spark Lightning-Fast Big Data Analysis](https://www.oreilly.com/library/view/learning-spark/9781449359034/)
<https://www.oreilly.com/library/view/learning-spark/9781449359034/>
- Coursera [Big Data Analysis with Scala and Spark](#)

- Spark provides an interactive environment where you can learn Spark quickly.
 - ▶ PySpark Shell: Use Python to interact with Spark, making it accessible for Python users.
 - ▶ Scala Shell: Use Scala, Spark's native language, for direct interaction with Spark.

- PySpark is the **Python API** for Apache Spark, allowing Python developers to harness the power of Spark's distributed computing capabilities.
 - ▶ By importing the `pyspark` library. Developers can then use its functions and classes to work with Apache Spark.
- An API, or Application Programming Interface, is a set of rules and protocols for building and interacting with software applications.
 - ▶ Web APIs: Accessed via HTTP, commonly used for web services (REST, SOAP).
 - ▶ Library APIs: Provided by software libraries, allowing applications to use predefined functions (e.g., Math libraries).
 - ▶ Operating System APIs: Allow applications to interact with the OS (e.g., Windows API).
 - ▶ Database APIs: Enable interaction with databases (e.g., SQL queries).

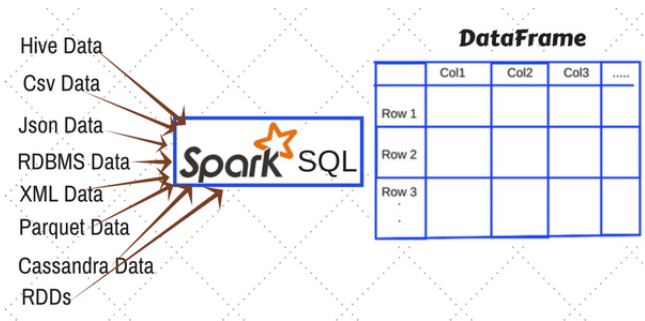
PySpark

- The Pyspark communicates with Spark using Py4J API



Data Frame

- Data Frame is a distributed collection of structured data. It is conceptually similar to a table in a relational database.
 - ▶ Data is distributed across multiple nodes in a cluster.
 - ▶ Data is organized into named columns.
 - ▶ Supports a wide range of data types (numeric, string, date, etc.).
- Spark SQL is one of the modules in the Spark ecosystem for structured data.
 - ▶ Allows users to query structured data using SQL as well as the DataFrame API.
 - ▶ Spark SQL provides different API's: DataFrame, RDD(Scala/Java), SQL



License



This work is licensed under a [Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-nc-sa/4.0/).