Data Engineering in the Cloud

Hybrid Transactional Analytical Processing

Xuemao Zhang
East Stroudsburg University

January 18, 2025

Outline

- In the labs, we Query Azure Cosmos DB with Apache Spark Pool and Serverless SQL Pools
- Lab 1: Querying Azure Cosmos DB with Apache Spark Pool
- Lab 2: Querying Azure Cosmos DB with Serverless SQL Pools

Prererquisites

- Prerequisites (10-15 minutes):
  - Create a datalake (storage account with hierarchical namespace enabled)
    - ★ To lower the cost, you may choose `Redundancy` as `LRS`
    - ★ Create a container in the storage account
  - Create an Azure Synapse workspace and
    - ★ create a Apache spark pool in the `Manage` tab
  - Create an Azure Cosmos DB database with a containter
    - ★ with `Analytical store` enabled and then create two key-value items
  - In the Synapse Studio, go to the `Manage` tab to add a new linked service for `Azure Cosmos DB for NoSQL`
  - Go to the `Data` tab of the Synapse Studio, click on + and choose "Integration dataset" to create a dataset.
    - ★ We do not need the Copy pipeline

Lab 1: Querying Azure Cosmos DB with Apache Spark Pool

- 1 Open "synapse studio", go to "Manage" tab, and add a new Cosmos DB `linked service`.

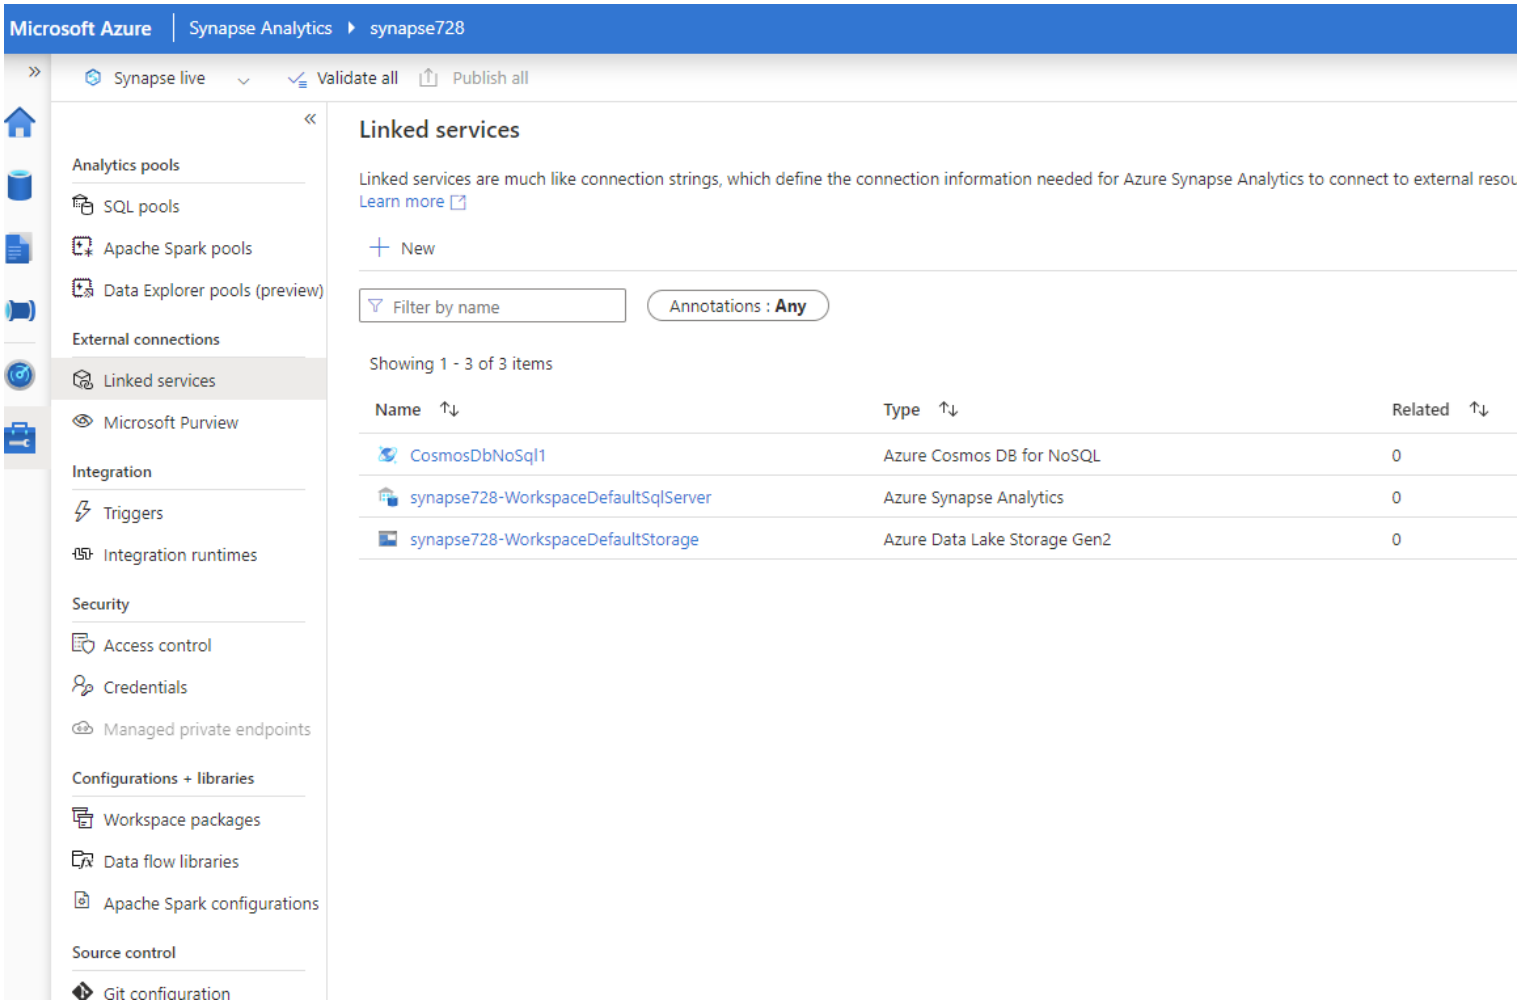

Lab 1: Querying Azure Cosmos DB with Apache Spark Pool

- 2 Go to “Data” tab, and choose “Linked”.

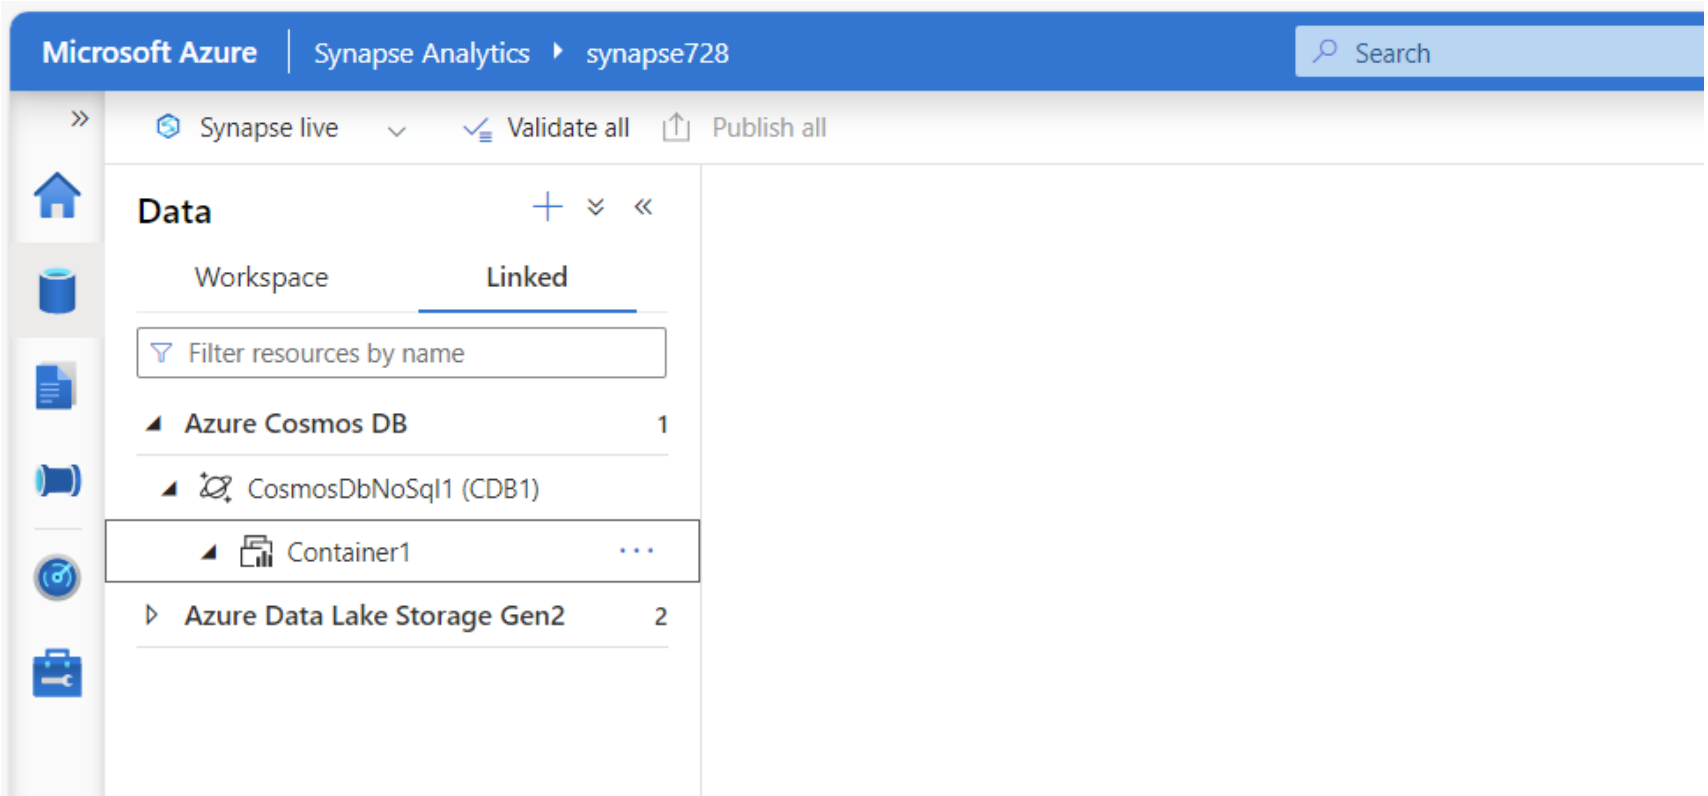

Lab 1: Querying Azure Cosmos DB with Apache Spark Pool

- 3 Expand the Azure Cosmos DB, choose the container "Container1" created in the Cosmos DB, click on three dots, choose "New notebook", and choose "Load to Dataframe".

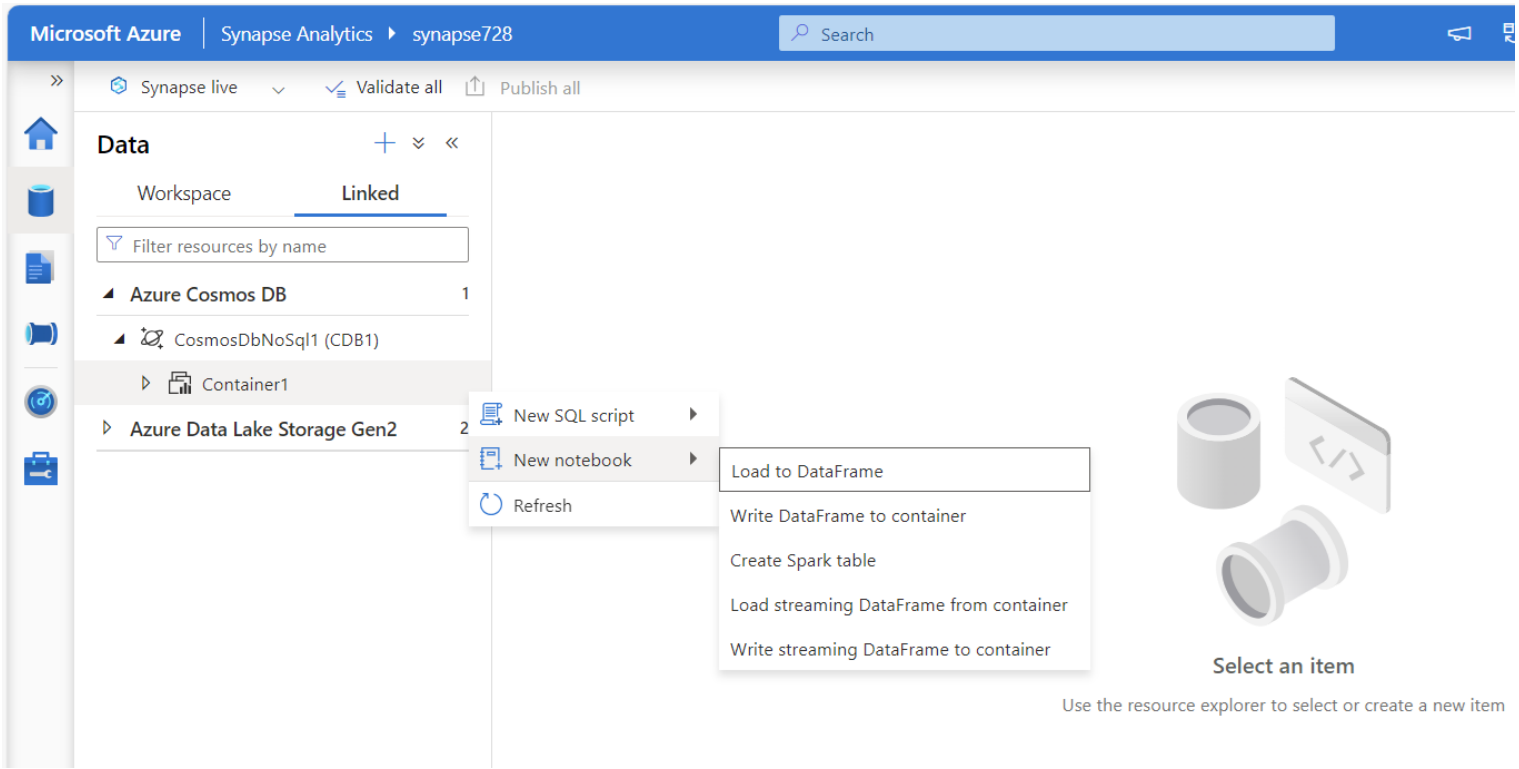

Lab 1: Querying Azure Cosmos DB with Apache Spark Pool

- 4 Attach the spark pool to the notebook.

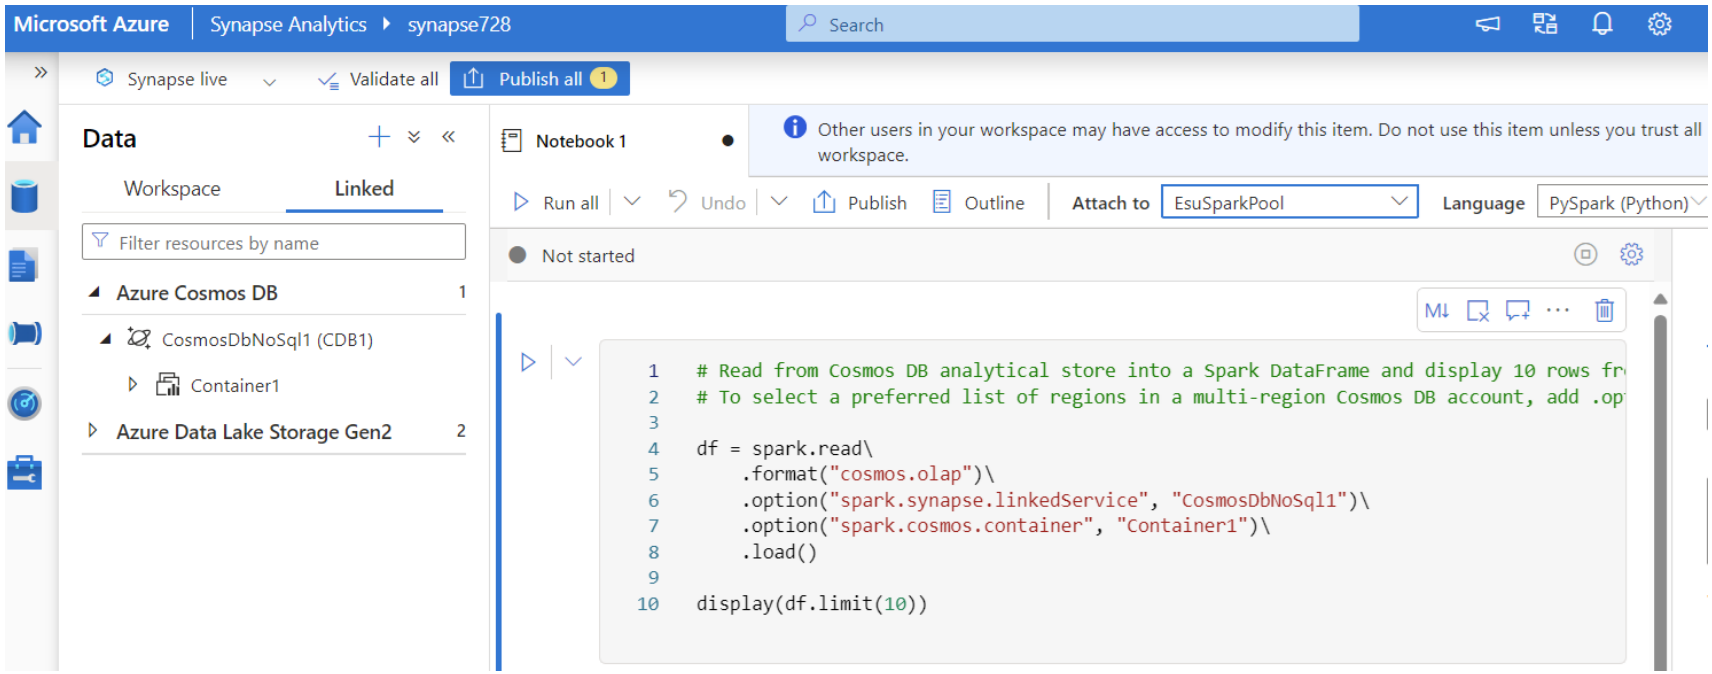

Lab 1: Querying Azure Cosmos DB with Apache Spark Pool

- 5 Click on "Run All". It takes some time to run.
  - If you see similar message, then update the configuration of the spark session

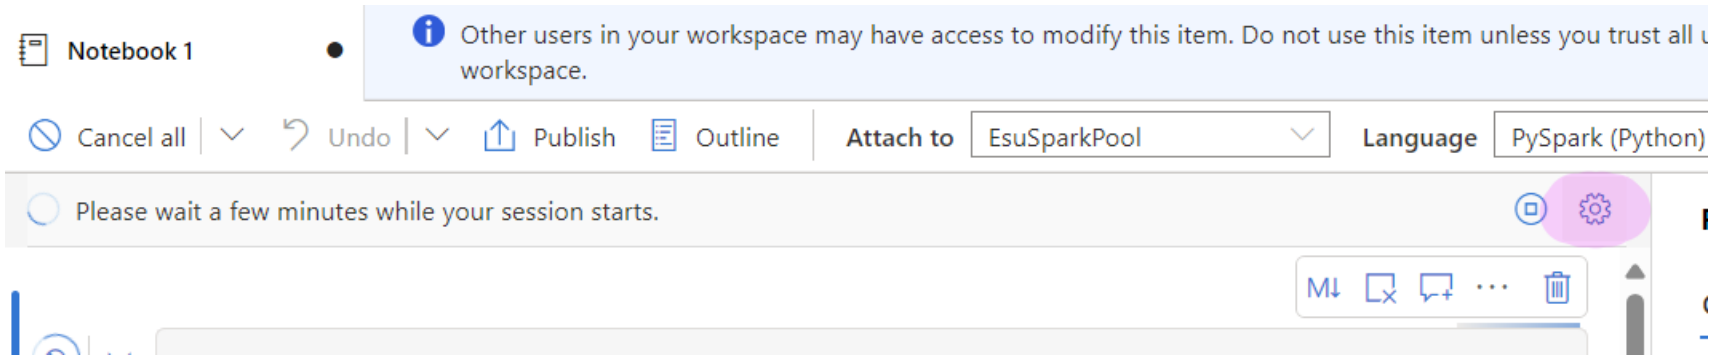


```
InvalidHttpRequestToLivy: Your Spark job requested 48 vcores.
However, the workspace has a 12 core limit. Try reducing the
numbers of vcores requested or increasing your vcore quota.
Quota can be increased using Azure Support request https://learn.microsoft.com/en-us
```

Lab 1: Querying Azure Cosmos DB with Apache Spark Pool

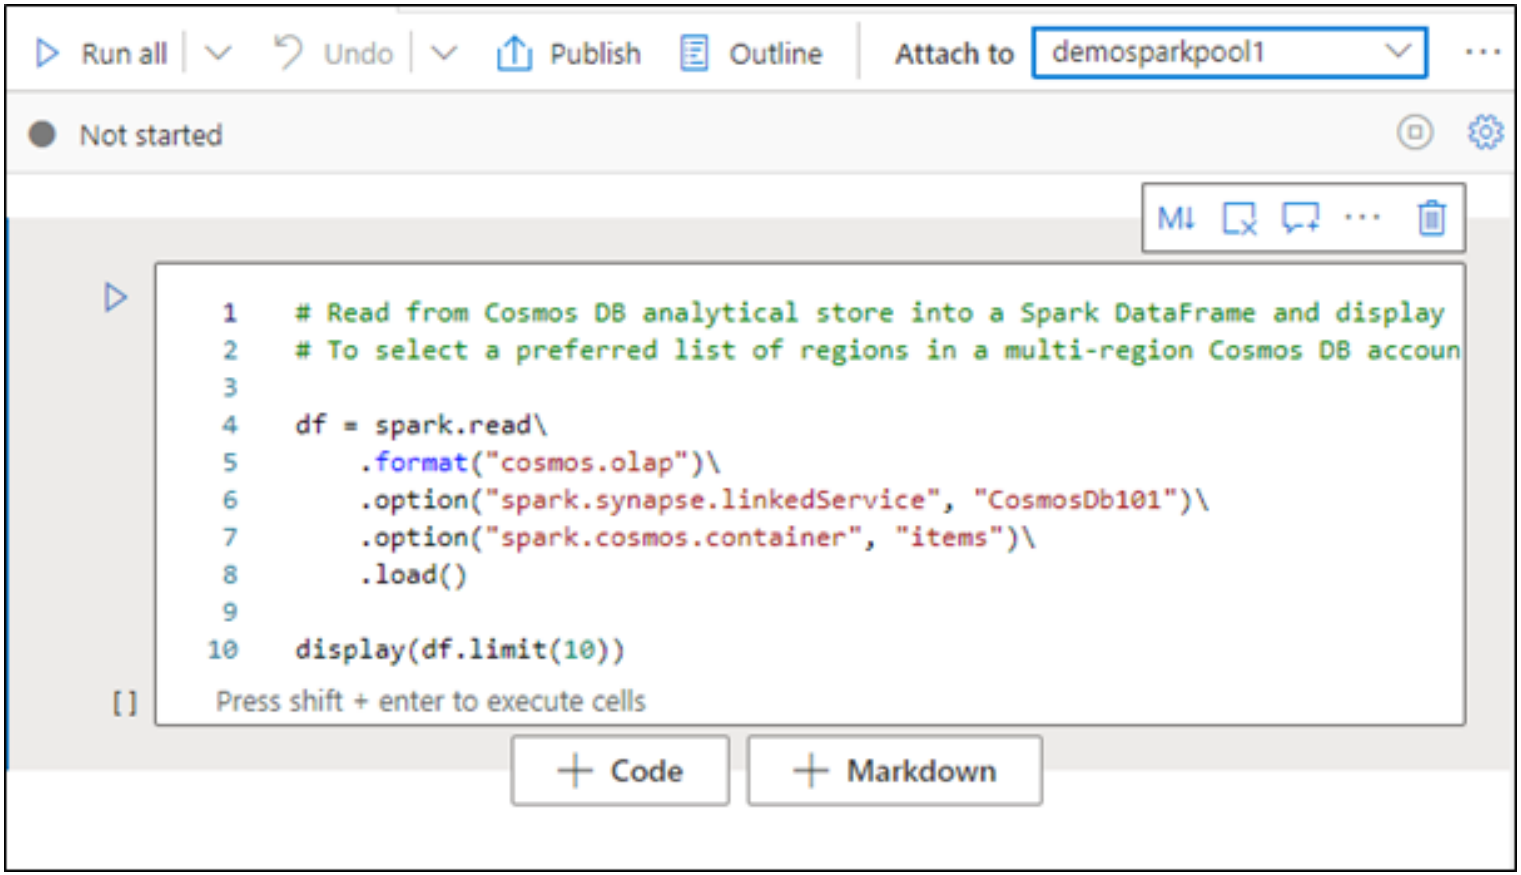

Lab 1: Querying Azure Cosmos DB with Apache Spark Pool

- 6 You must be able to see the output as given below:

```python
# Read from Cosmos DB analytical store into a Spark DataFrame and display 10 rows from the DataFrame
# To select a preferred list of regions in a multi-region Cosmos DB account, add .option("spark.cosmos.preferredRegions", "<Region1>,<Region2>")

df = spark.read\
    .format("cosmos.olap")\
    .option("spark.synapse.linkedService", "CosmosDbNoSql1")\
    .option("spark.cosmos.container", "Container1")\
    .load()

display(df.limit(10))
```

[1] ✓ 3 min 28 sec - Apache Spark session started in 3 min 2 sec 344 ms. Command executed in 25 sec 972 ms by niuyuxia728 on 2:59:48 PM, 7/28/24

Job execution Succeeded **Spark** 1 executors 4 cores View in monitoring Open Spark UI

View Table Chart Export results

_rid	_ts	id	category	description	_etag
2J4oAMexroYBAAAAAAAAAA==	1722190945	1	fruits	vrious fruits	"06007145-0(
2J4oAMexroYCAAAAAAAAAA==	1722190981	2	vegetables	various vegetables	"06005846-0(

Lab 1: Querying Azure Cosmos DB with Apache Spark Pool

- 7 You can remove the unwanted column using the below command:

```
unwanted_cols = {'_rid','_ts','_etag'}
cols = list(set(df.columns) - unwanted_cols)
df2 = df.select(cols)

display(df2.limit(10))
```

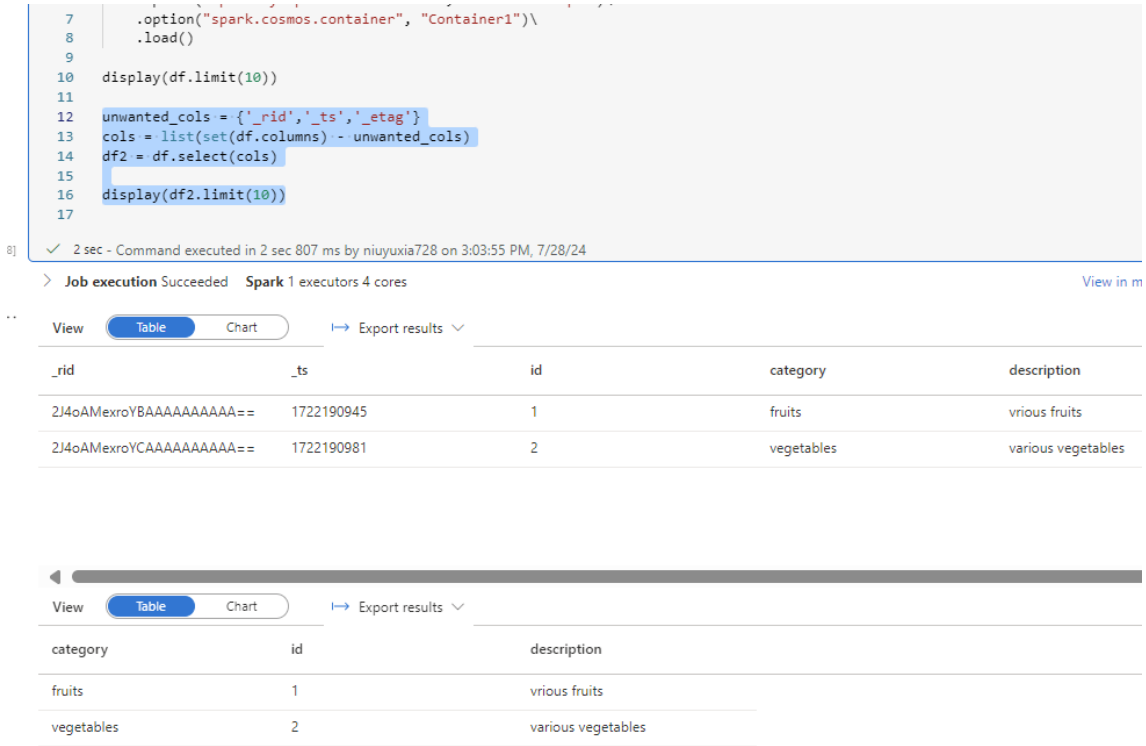

```
7       .option("spark.cosmos.container", "Container1")\
8       .load()
9
10  display(df.limit(10))
11
12  unwanted_cols = {'_rid','_ts','_etag'}
13  cols = list(set(df.columns) - unwanted_cols)
14  df2 = df.select(cols)
15
16  display(df2.limit(10))
17
```

8] ✓ 2 sec - Command executed in 2 sec 807 ms by niuyuxia728 on 3:03:55 PM, 7/28/24

Job execution Succeeded Spark 1 executors 4 cores View in m

View Table Chart Export results

_rid	_ts	id	category	description
2J4oAMexroYBAAAAAAAAAA==	1722190945	1	fruits	vrious fruits
2J4oAMexroYCAAAAAAAAAA==	1722190981	2	vegetables	various vegetables

View Table Chart Export results

category	id	description
fruits	1	vrious fruits
vegetables	2	various vegetables

Lab 2: Querying Azure Cosmos DB with Serverless SQL Pools

- 1 Go to "Data" in the Synapse Studio and right click `Container1` to add an SQL script.

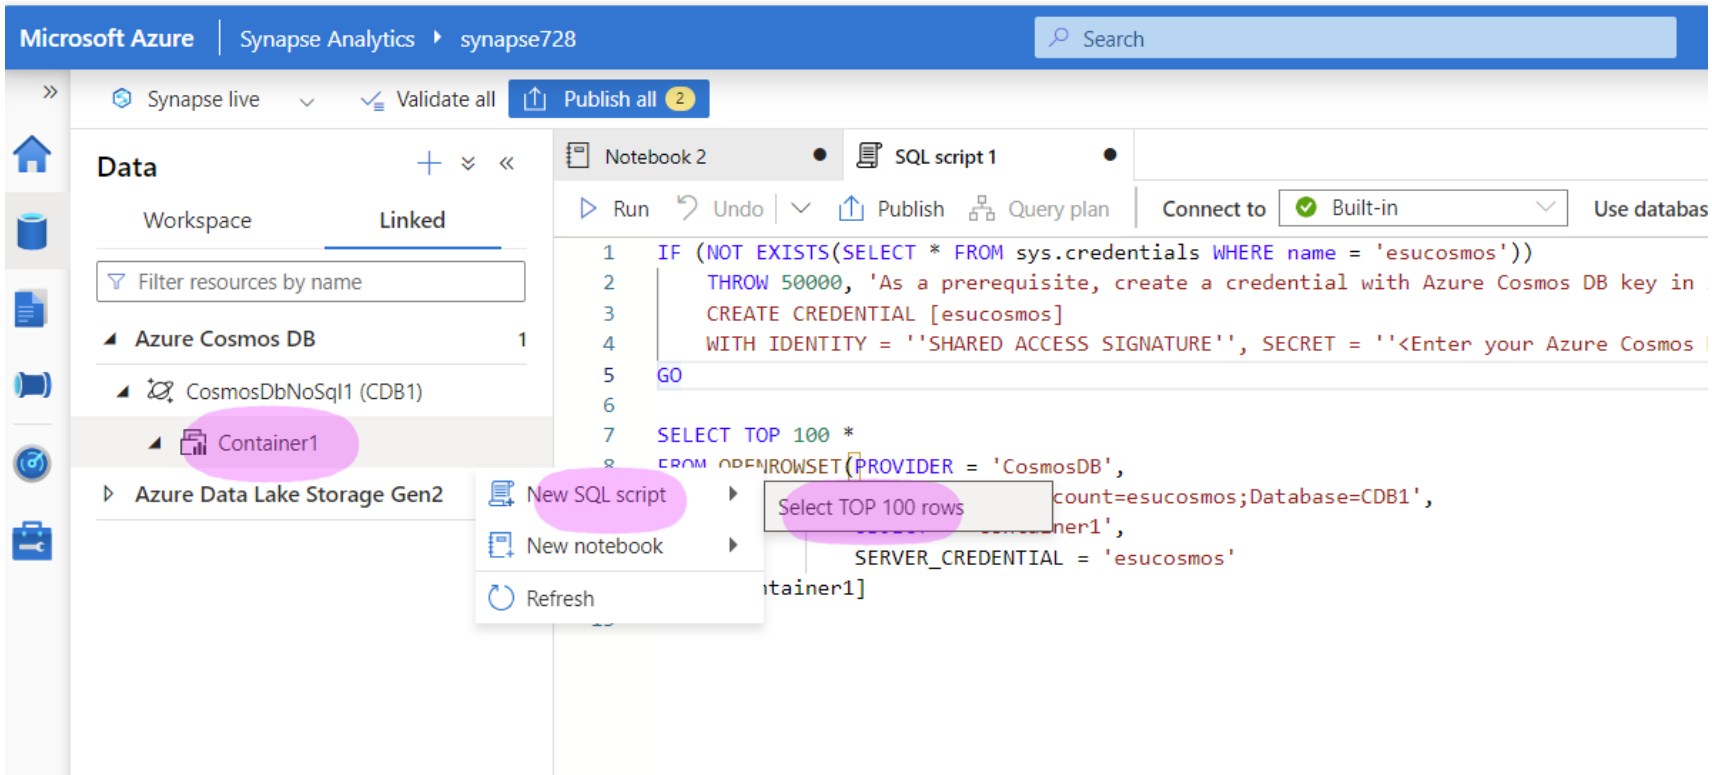

Lab 2: Querying Azure Cosmos DB with Serverless SQL Pools

- 2 Run the first block query see if a credential is needed.

```
IF (NOT EXISTS(SELECT * FROM sys.credentials WHERE name = 'esucosmos'))
    THROW 50000, 'As a prerequisite, create a credential with Azure Cosmos DB key in SECRET option:
    CREATE CREDENTIAL [esucosmos]
    WITH IDENTITY = ''SHARED ACCESS SIGNATURE'', SECRET = ''a4WXnMntqa0yJfMR4D84rcfNgDCDR33dfo70PMeum8nOvNeshAEqkD0TdyycfyHUFyrDrQUAbW
GO
```

- 3 Create a credential first

```
CREATE CREDENTIAL [esucosmos]
WITH IDENTITY = 'SHARED ACCESS SIGNATURE',
SECRET = 'a4WXnMntqa0yJfMR4D84rcfNgDCDR33dfo70
PMeum8nOvNeshAEqkD0TdyycfyHUFyrDrQUAbWUNACDb8oZYVg==';
```

Lab 2: Querying Azure Cosmos DB with Serverless SQL Pools

- 4 Run the query:

```
SELECT TOP 100 *
FROM OPENROWSET( PROVIDER = 'CosmosDB',
                 CONNECTION = 'Account=esucosmos;Database=CDB1',
                 OBJECT = 'Container1',
                 SERVER_CREDENTIAL = 'esucosmos'
) AS [Container1]
```

Lab 2: Querying Azure Cosmos DB with Serverless SQL Pools

```
10
11  SELECT TOP 100 *
12  FROM OPENROWSET(PROVIDER = 'CosmosDB',
13                  CONNECTION = 'Account=esucosmos;Database=CDB1',
14                  OBJECT = 'Container1',
15                  SERVER_CREDENTIAL = 'esucosmos'
16  ) AS [Container1]
17
```

Results Messages

View Table Chart ⟼ Export results

Search

description	category	_rid	_etag	_ts	id
vrious fruits	fruits	2J4oAMexroYBAAA...	"06007145-000...	1722190945	1
various vegetables	vegetables	2J4oAMexroYCAAA...	"06005846-000...	1722190981	2

Lab 2: Querying Azure Cosmos DB with Serverless SQL Pools

- 5 Now let's creat a database in Azure Synapse

```
-- Create the Profiles database if it does not exist
USE master;
GO
IF DB_ID(N'Profiles') IS NULL
BEGIN
    CREATE DATABASE Profiles;
END
GO

-- Switch to the Profiles database
USE Profiles;
GO

-- Drop the items view if it exists
DROP VIEW IF EXISTS items;
GO
```

Lab 2: Querying Azure Cosmos DB with Serverless SQL Pools

```
-- Create the items view to access Cosmos DB
CREATE VIEW items
AS
SELECT
    *
FROM OPENROWSET(
PROVIDER = 'CosmosDB',
CONNECTION = 'Account=esucosmos;Database=CDB1',
OBJECT = 'Container1',
SERVER_CREDENTIAL = 'esucosmos'
)
WITH (
    id varchar(1000),
    category varchar(50),
    description varchar(max)
) AS profiles;
GO
```

Lab 2: Querying Azure Cosmos DB with Serverless SQL Pools

- 6 Check the new database

```
USE Profiles;
GO

SELECT TOP 10 *
FROM items;
```

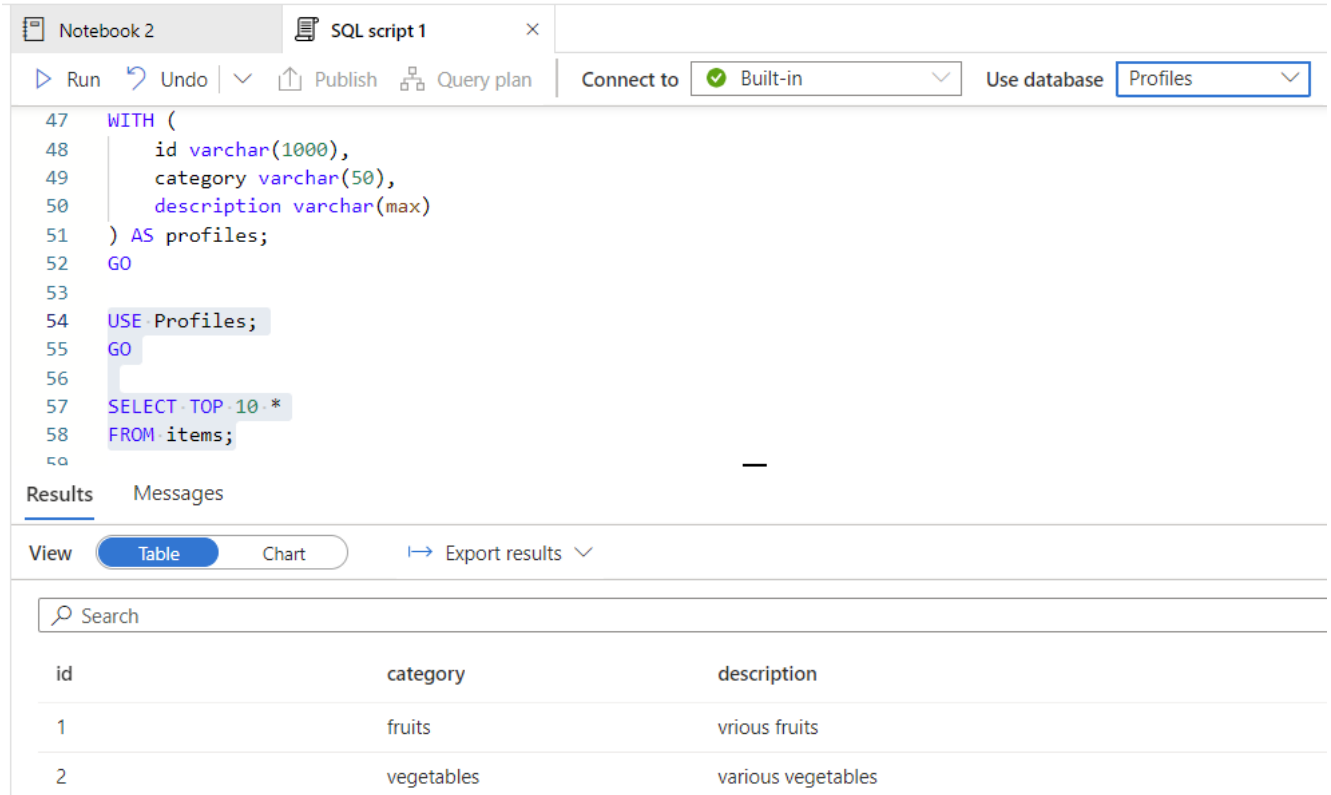

Lab 2: Querying Azure Cosmos DB with Serverless SQL Pools

- 7 Check the new database: Refresh the Synapse Studio, then go to "Data" tab, expand the "Profiles (SQL)" and under "Views" there must be a view created.

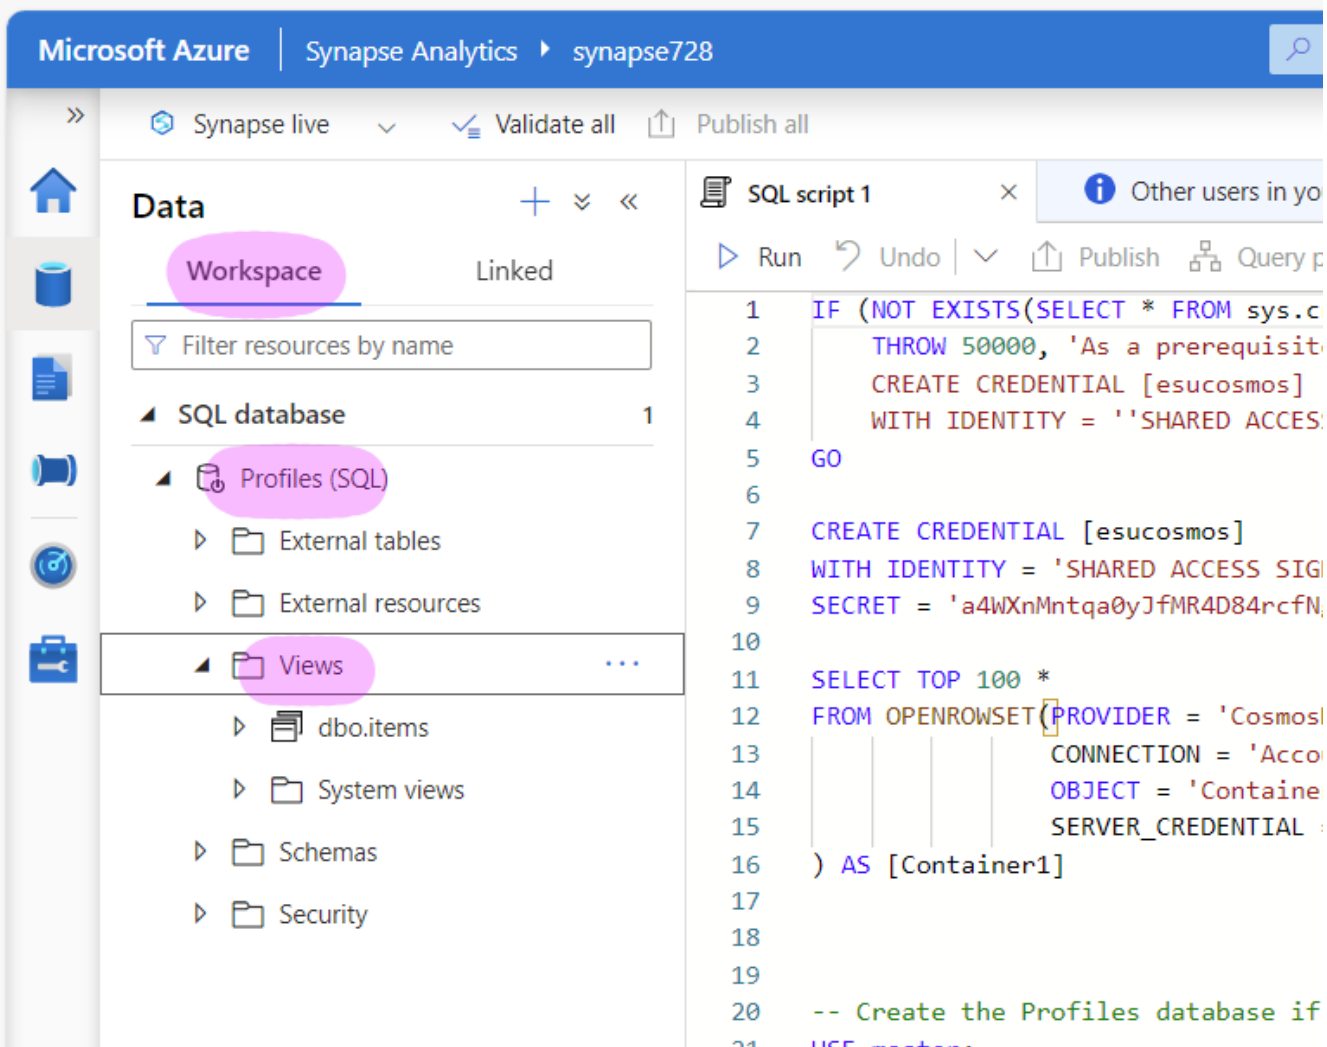

Lab 2: Querying Azure Cosmos DB with Serverless SQL Pools

- 8 Click on "dbo.items", choose "New SQL script", click "Select Top 100 rows" and you will be able to see the auto generated query. Run the query.

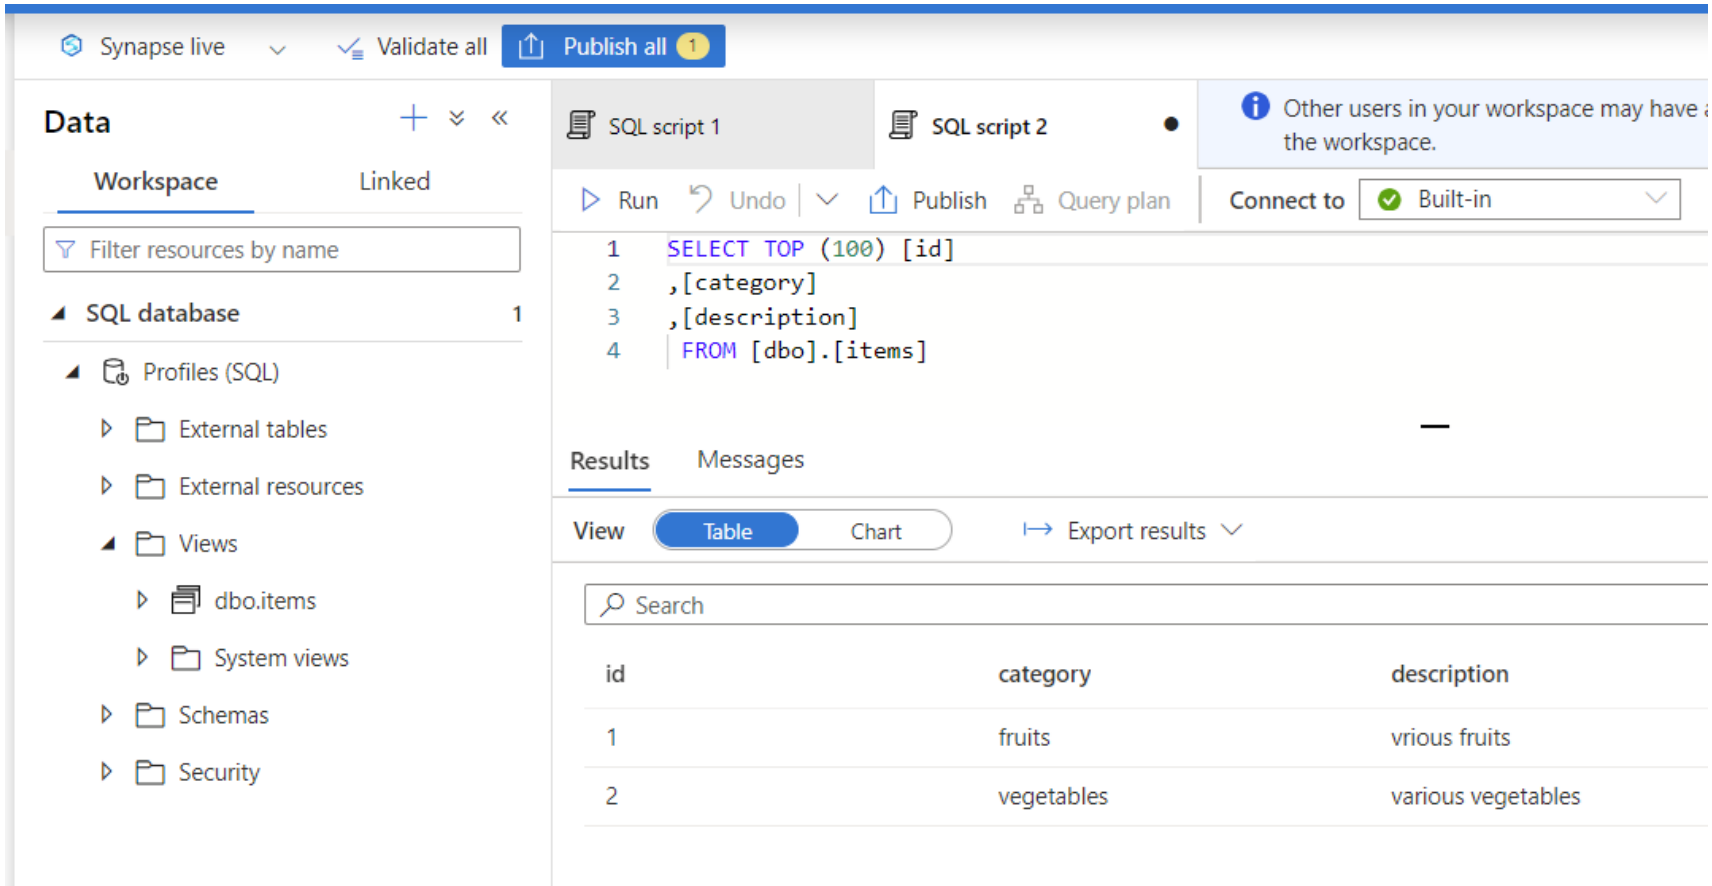

License